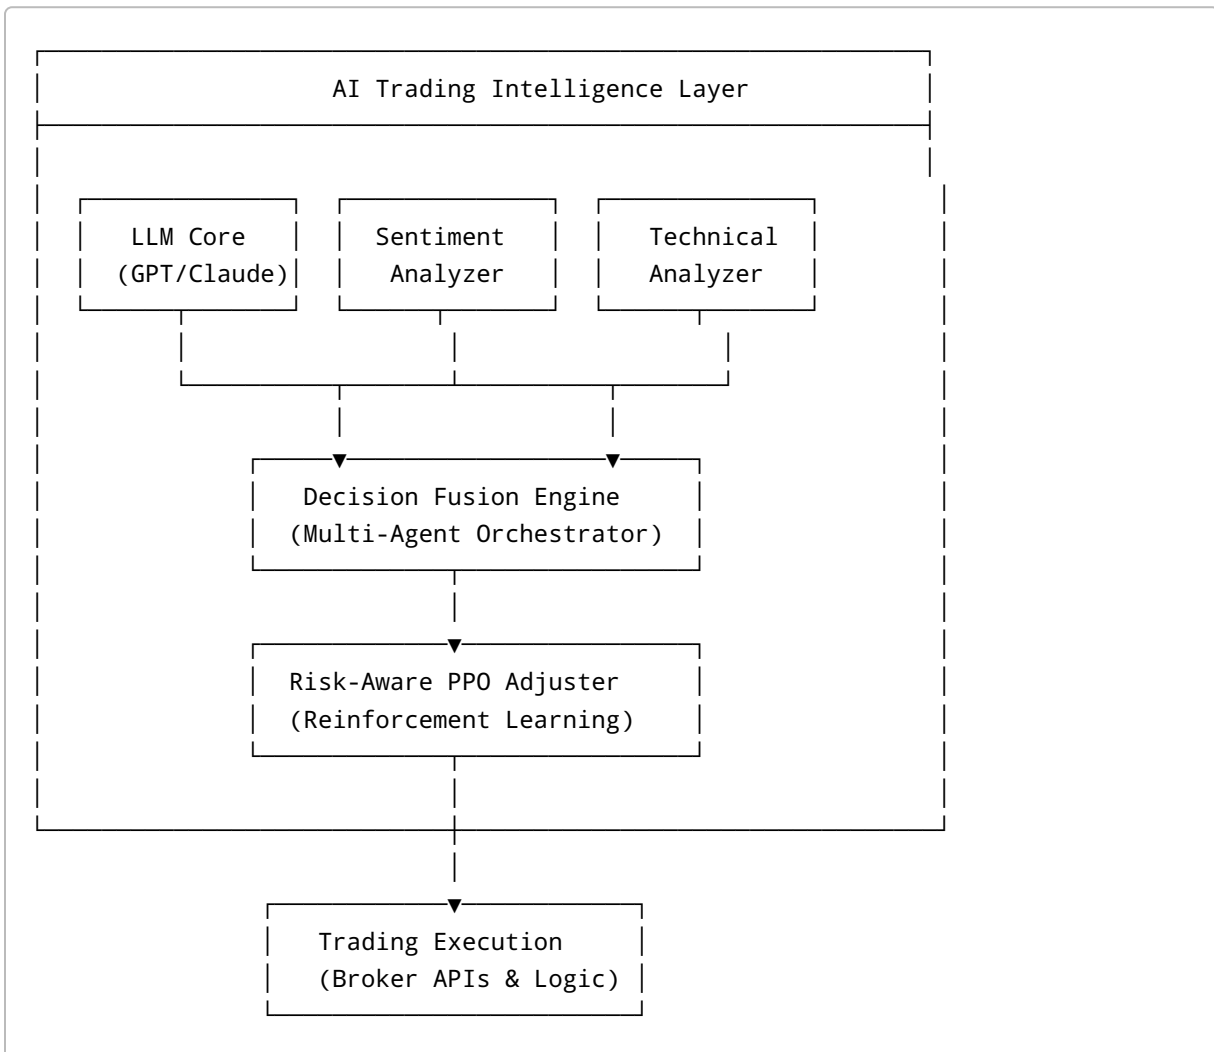


AI-Enhanced Crypto Trading Bot: Technical Design & Implementation Plan

Architecture Overview

The AI-enhanced trading bot is structured as a modular, multi-layer system combining traditional algorithmic trading logic with advanced AI components. The design emphasizes **modularity**, **asynchronous processing**, and **extensibility**. At a high level, the bot's architecture can be visualized as follows:



Component Roles: In this architecture, **market data and external information flow into specialized modules** (Technical Analyzer, Sentiment Analyzer, LLM Predictor). These feed into a **Decision Fusion Engine**

(a multi-agent system) which combines insights and makes a trading decision ¹. The decision is then passed through a **Risk-Aware Adjuster** (PPO-based reinforcement learning module) that tweaks the action based on risk metrics (e.g. VaR/CVaR) ². Finally, a unified **Execution & Order Routing** layer interfaces with broker/exchange APIs to execute trades. Surrounding this core are cross-cutting systems for **configuration**, **data input**, and **monitoring**.

Asynchronous, Two-Tier Decision Making: To handle real-time requirements, the design supports a dual-path decision process. A fast rule-based path can act immediately on obvious signals, while an AI-assisted path runs in parallel for deeper analysis of less clear situations ³. This asynchronous approach ensures that time-sensitive opportunities are not missed due to LLM latency ⁴.

Codebase Scaffolding: The project is organized for clarity and extensibility. A possible directory/module layout is:

- `core/` – Core trading logic and strategies (strategy definitions, Technical Analyzer, StrategyContext, etc.)
- `ai/` – AI modules (LLM Manager, AITrendPredictor, multi-agent Decision Engine, PPO Adjuster, Sentiment Engine)
- `risk/` – Risk management (RiskManager, VaR/CVaR calculations, position sizing logic)
- `execution/` – Broker/exchange adapters and Execution Orchestrator
- `data/` – Data handling (market data connectors, data quality checks, backtesting utilities)
- `config/` – Configuration files and Pydantic models
- `interfaces/` – Interface layers (FastAPI web server, CLI via Typer, optional GUI)
- `tests/` – Unit and integration tests for each module

This structure ensures each concern is in a dedicated module, making it easy to maintain or replace components (for example, adding a new exchange adapter or swapping out the LLM provider).

Technology Stack

Programming Language & Base Platform: The bot is built with **Python 3.11+** for its rich ecosystem of data science libraries and async support. Python's latest version ensures improved performance and syntax features (e.g. `asyncio` improvements) crucial for real-time trading.

Backend Framework: **FastAPI** is recommended for building a RESTful API service for the bot. FastAPI allows asynchronous request handling (matching our async design) and can serve a web-based dashboard or receive trading commands remotely. It also integrates Pydantic for data validation (useful for config and input validation).

Data Processing: **PyArrow** is included for efficient in-memory data handling and interchange. Market data can be large (multi-symbol, high-frequency); using PyArrow helps with zero-copy data sharing between components and can serialize data frames to Apache Arrow format for fast transport or persistence.

CLI Interface: **Typer** (a library built on Click) is used to implement a **command-line interface**. The CLI will support operations like starting the bot, running backtests, and switching configuration profiles. Typer

offers a simple way to define commands, options, and help text, making the bot easily controllable via terminal.

User Interface (Optional): For a rich user interface, two options are considered: - A **ReactJS** front-end web dashboard (communicating with FastAPI backend). This would display real-time metrics, open positions, P/L charts, and allow strategy monitoring and configuration in a browser. - A **PySide6** (Qt) desktop application for a local GUI, if a native app is preferred. This can reuse the same API endpoints or Python modules for data.

These UI components are optional and can be introduced once the core bot logic is stable. Initially, monitoring can be done via CLI and logs.

Libraries & Tools: Additional libraries include: - **TA-Lib** (or `ta` pandas-based library) for technical indicator calculations (RSI, MACD, Bollinger Bands, etc.). - **CCXT** (or similar) for multi-exchange API support, providing a unified interface to various crypto exchanges for data and trading. - **OpenAI/Anthropic SDKs** or HTTP clients for LLM integration (if using external APIs), and possibly `llama.cpp` or HuggingFace Transformers for local models. - **Stable-Baselines3** or custom RL implementation for PPO training, if using a high-level library to implement reinforcement learning. - **Prometheus client** for metrics, **python-telegram-bot** or HTTP API for Telegram messaging, and standard Python logging (with JSON or structured logs) for observability.

This stack ensures that the system is built on proven, well-supported technologies and can be deployed locally or scaled to cloud infrastructure (Docker/Kubernetes for GCP in production).

Modular Design

To maximize extensibility, the bot is broken into layers/modules each handling a specific domain of responsibility. All modules communicate via well-defined interfaces (e.g., data classes or async API calls), enabling independent development and testing.

LLM Manager (Multi-Provider AI Interface)

The **LLM Manager** module provides a unified interface to multiple Large Language Model providers (OpenAI, Anthropic Claude, local models, etc.). It encapsulates the complexity of calling different APIs behind a common asynchronous `generate(prompt, system_message)` method. Key features:

- **Provider Abstraction & Fallback Chain:** The manager can switch between providers based on availability or cost. For example, if the primary (OpenAI) fails or is too slow, it automatically falls back to another (Anthropic, local) ⁵ ⁶ ⁷ . This prevents a single point of failure and improves resilience.
- **Model Selection & Configurability:** The specific model used (e.g. GPT-4 vs GPT-3.5, Claude 1 vs 2) is configurable in the YAML settings, and the manager will initialize the appropriate client accordingly ⁸ ⁹ . New providers can be added by implementing a small adapter for the LLM API.
- **Cost Tracking & Limits:** The manager tracks usage and cost per provider. A monthly budget can be set (e.g. \$100/month) and the system will avoid exceeding this by degrading to offline strategies if necessary ⁷ .

- **Caching Responses:** Repeated LLM queries (e.g. asking about the same symbol multiple times in an hour) are cached for a configurable TTL to save time and cost ⁷.
- **Safe Defaults:** The LLM Manager includes safeguards such as a **circuit breaker** to disable a provider after repeated failures and a **response validator** to check that outputs are valid and not hallucinated ¹⁰ ¹¹. It also enforces structured output (e.g. JSON-only responses as per a strict system prompt) to ensure downstream components can parse the AI predictions reliably.

By centralizing these features, any part of the bot that needs language model intelligence (trend prediction, sentiment analysis, etc.) can call the LLM Manager without worrying about provider-specific details.

Sentiment Engine (Multi-Source Sentiment Analyzer)

The **Sentiment Engine** gathers and aggregates market sentiment from various sources in real-time. Rather than relying on a single expensive API (like the Twitter API alone), it uses a combination of **free or low-cost data sources for robustness and cost-efficiency** ¹²:

- **Social Media:** Using Reddit (via PRAW or Pushshift) and Twitter API (limited use or filtered via third-party services) to gauge retail sentiment (mentions, positive/negative ratios, trending topics).
- **News Feeds:** Parsing crypto news RSS feeds (e.g. CoinTelegraph, Decrypt) to perform headline sentiment analysis ¹³. This can be done with a lightweight NLP model or keyword scoring to categorize news as positive, neutral, or negative.
- **Market Sentiment Indices:** Integrating indices like the Fear & Greed Index (often available for free daily) ¹⁴.
- **On-chain/Exchange Data:** Leverage exchange-provided sentiment indicators like funding rates and long/short ratios from futures markets ¹⁵, which often signal trader sentiment (e.g., consistently high funding rates may indicate bullish sentiment).

The Sentiment Engine runs asynchronously, pulling data from these sources on a schedule (and caching where appropriate to avoid rate limits). It produces a **combined sentiment score** or structured data object that includes sub-components (social sentiment, news sentiment, fear/greed, etc.). An example approach is to compute a weighted average of available sentiment metrics ¹⁶, putting more weight on high-signal sources (like exchange funding data) and gracefully handling missing data if an API is down.

This module will provide functions like `analyze(symbol) -> SentimentData` which can be used by strategy logic or the AITrendPredictor. By designing it as a separate service, it's easy to update (e.g., add a new social platform) or disable expensive sources, and to reuse sentiment data for multiple strategies or symbols simultaneously.

Technical Analyzer (Indicator Computation & Pattern Detection)

The **Technical Analyzer** is responsible for all quantitative market data analysis. It ingests price and volume data for multiple symbols and timeframes (through exchange APIs or historical data sources) and computes a suite of technical indicators and features. Key elements include:

- **Indicator Library:** Using TA-Lib or a pandas-based library, the analyzer computes common indicators like **RSI**, **MACD**, **Bollinger Bands**, **Moving Averages (MA20, MA50, MA200)**, **ATR (Average True Range)**, etc. ¹⁷. These indicators provide insight into momentum, trend, volatility, and mean reversion signals for each symbol.

- **Pattern Recognition:** It can detect chart patterns or candlestick formations (if needed) and identify support/resistance levels. For example, recent local minima/maxima can be flagged as key support/resistance, and volatility measures can help judge significance.
- **Multi-Timeframe Support:** The analyzer can fetch data for multiple timeframes (e.g. 1h, 4h, 1d) in parallel and compute indicators for each. This is important for strategies that consider a bigger trend context even for short-term trades.
- **Data Preprocessing & Validation:** Before computing indicators, the Technical Analyzer uses a **DataQualityMonitor** to validate incoming data (checking for gaps, outliers, stale timestamps, etc.)¹⁸ ¹⁹. This ensures that bad data (e.g., a faulty price spike or missing candles) doesn't lead to false signals. Any anomalies can trigger either a data refresh or a flag for the risk manager to halt trading on that symbol until resolved.

The output of this module is a structured **TechnicalSignals** object per symbol/timeframe, containing the latest values of all relevant indicators and possibly simple interpretations (e.g., “RSI overbought/oversold” flags, moving average crossover signals). This technical snapshot is used by both traditional rule-based strategies and the AI components (it can be embedded into LLM prompts or fed into the PPO model).

AITrendPredictor (LLM-Powered Multi-Timeframe Forecasting)

The **AI Trend Predictor** is a higher-level module that leverages the LLM Manager, Sentiment Engine, and Technical Analyzer together to produce an overall market outlook. Its goal is to analyze multiple data modalities and forecast price trends across several time horizons, providing a rich context for decision-making.

Functionality: For a given symbol (or set of symbols), the AITrendPredictor will gather the latest **multi-timeframe data** (e.g. 1h, 4h, 1d, 1w) and for each timeframe perform the following steps:

1. **Gather Data:** Fetch recent market data (candles, stats) for that timeframe, compute technical signals, and retrieve current sentiment data and recent news headlines relevant to the asset²⁰ ²¹.
2. **Build Prompt:** Construct a comprehensive prompt that summarizes the symbol's state – including technical indicator readings, trend metrics, sentiment summary, and key news items. The prompt is structured to give the LLM a complete picture (for example, listing prices, volatility, indicator values, sentiment scores, etc., followed by a clear instruction to provide a trend prediction in JSON format)²² ¹⁷. This prompt engineering is crucial to ensure the LLM's output is grounded in the data provided and follows a strict format.
3. **LLM Prediction:** Send the prompt to the LLM Manager to get a prediction. The system prompt ensures the LLM acts as a trading analyst and returns a **JSON** with fields like predicted direction (bullish/bearish), confidence level, expected price movement, reasoning bullets, etc. according to a predefined schema²³ ²⁴.
4. **Post-process LLM Output:** Parse the LLM's JSON output into a Python object (e.g., a `TrendPrediction` dataclass). Validate that the content makes sense (via the LLM Manager's response validator, which checks for hallucinations or format issues¹¹).
5. **Risk Adjustment (PPO):** Before finalizing the prediction for that timeframe, pass it through the PPO-based risk adjuster (described below). This will adjust certain aspects like recommended position size or confidence based on quantitative risk considerations (for example, reducing confidence if predicted move contradicts risk metrics or if market volatility is high).
6. **Aggregate Multi-Timeframe Insights:** Combine the predictions from each timeframe into an overall view²⁵ ²⁶. For instance, the engine might require alignment between short-term and long-

term outlooks before high confidence, or it may derive a composite signal (e.g., if 1h and 4h are bullish but 1d is neutral, overall moderate bullish outlook). This aggregation logic yields a final **AI-driven signal** for the symbol.

By using multi-modal data (price, technicals, sentiment, news) and multi-horizon analysis, the AITrendPredictor aims to produce more nuanced signals than any single indicator. Research has suggested LLM-enhanced strategies can improve prediction accuracy by 15-30% ²⁷ ²⁸, although caution is taken not to overestimate performance. All predictions include a measure of confidence and explicit reasoning, which can be logged or presented to users for transparency.

PPO-Based Adjuster (Risk-Aware Reinforcement Learning Module)

The **PPO-Based Adjuster** is a reinforcement learning component that fine-tunes trade decisions or parameters with a focus on risk management. It implements a variant of Proximal Policy Optimization (PPO) to adjust the outputs of the AI or strategy before execution, in order to optimize risk-adjusted returns.

Role in the System: This module takes initial trade plans (from either the AITrendPredictor or a human-defined strategy) and adjusts elements like position size, entry/exit thresholds, or confidence based on learned policies. It acts as an automated risk advisor that *learns from historical data and simulations* to improve outcomes (e.g., avoid trades when risk is disproportionate to expected return).

Key design aspects:

- **Risk-State Representation:** The adjuster maintains a state representation that includes current risk metrics (VaR, CVaR, volatility), portfolio exposures, recent performance, etc. For example, it calculates daily Value-at-Risk at 95% confidence and Conditional VaR ²⁹ for the portfolio using a VaRCalculator and CVaRCalculator. These values, along with metrics like current leverage or correlation of the asset with the portfolio, form the state input to the PPO model.
- **Policy Actions:** The PPO agent's action might be a scaling factor or adjustment recommendation. For instance, if the LLM suggests a 5% allocation long but current VaR usage is high, the PPO policy might scale it down (e.g., cut position by 50% if VaR is near limit) ³⁰. Similarly, it could adjust the confidence or override a trade (to "no trade") if risk is extreme.
- **Training:** The PPO model can be trained in the **backtesting engine** (offline mode) where it simulates many trading episodes using historical data. The reward function would be formulated to encourage higher risk-adjusted returns (e.g., Sharpe ratio improvements, penalizing drawdowns or VaR limit breaches). Over time, the PPO learns when to throttle back trades (or even increase size when risk is low and confidence is high) better than static rules.
- **Integration:** In live operation, the adjuster runs quickly to refine each decision. It can be implemented as an *async call that wraps around the decision-making process*, adjusting outputs before they reach the Execution layer. Because it's trained to be risk-aware, it introduces a safety net: for example, **if volatility regime is high or portfolio drawdown is worsening, the adjuster will reduce new position sizes or recommend skipping trades** (aligning with rules like "no leverage in high volatility" from the RiskManager) ³¹.

The PPO adjuster works in tandem with the rule-based risk management. While RiskManager (below) applies hard constraints and checks, the PPO adjuster provides a softer, continuous optimization of risk vs reward. This dual approach ensures both **hard safety limits and adaptive adjustments** are in place. By using reinforcement learning, the system can adapt to market behavior changes over time (for instance, learning to avoid trades ahead of certain economic news that historically led to whipsaws, if such patterns are in the training data). Importantly, this component is optional at runtime – if disabled or not fully trained, the system will fall back to static risk rules (safe default behavior).

RiskManager (Dynamic Sizing & Risk Controls)

The **RiskManager** module handles traditional risk management functions essential for safe trading operations. It enforces **safe defaults** and risk limits at all times, independent of any AI decisions. Major responsibilities include:

- **Pre-Trade Checks:** Before any trade is executed, the RiskManager performs a series of checks on the trade plan and current market conditions ³² ³³. For example, it will simulate the expected slippage by looking at order book depth for the trade size; if slippage would be too high (e.g. >0.5%), it may reject or downsize the order ³⁴. It also checks liquidity (e.g. not trading a position larger than 1% of daily volume) ³⁵, correlation (ensure new trades don't overly concentrate portfolio in highly correlated assets), and even time-of-day or market regime constraints (avoid trading during known illiquid periods or when a "flash crash" scenario is detected) ³¹ ³⁶.
- **Dynamic Position Sizing:** Using inputs like asset volatility (ATR), account equity, and risk tolerance, the RiskManager computes position sizes for each trade. For example, it might target a fixed percentage of portfolio value at risk per trade (e.g. 1% VaR) and adjust quantity accordingly. If the AI suggests a trade with high confidence but the asset is extremely volatile, the RiskManager will cap the size to limit potential loss. It also factors in open positions (if any) to ensure overall exposure stays within limits (e.g., maximum leverage or sector exposure).
- **Portfolio Drawdown Management:** The RiskManager monitors overall portfolio performance and can implement **kill-switches or cooldowns**. If cumulative drawdown exceeds a threshold (say 10%), the system might pause trading or only allow reduced-size positions until recovery. This is a safety feature to prevent a series of losses from spiraling.
- **Emergency Stop System:** A special sub-component monitors for catastrophic conditions – such as exchange connectivity issues, API failures, or drastic market moves (flash crashes). If triggered, it will **cancel all orders, close positions, and halt the bot** immediately ³⁷ ³⁸, while notifying the operators. This ensures the bot can gracefully shut down or step aside when things go seriously wrong (rather than exacerbating the problem).
- **Post-Trade Tracking:** After trades, the RiskManager logs outcomes and updates metrics like current VaR usage, realized volatility, win/loss streaks, etc. This data is fed back into both the PPO adjuster (for learning) and to the monitoring systems.

The RiskManager is largely rule-based and deterministic, serving as the **last line of defense**. Even if the AI/LLM components make an aggressive suggestion, the RiskManager can veto or modify it if it violates risk policies. These policies are configurable via the YAML config (for example, max allowed leverage, max position per asset, VaR limits, trading schedule, etc.). By separating this logic, one can adjust risk parameters without touching the strategy code, and ensure any strategy (AI or otherwise) adheres to the same risk constraints.

StrategyContext and Multi-Strategy Support

The bot is designed to support **multiple trading strategies and symbols concurrently**. The **StrategyContext** abstraction helps manage this complexity. Each StrategyContext can be thought of as an instance of a strategy running on a given market (or set of markets), maintaining its own state and parameters.

- **Encapsulation of Strategy Logic:** A StrategyContext ties together the data feeds, analyzers, and decision logic for a particular strategy. For example, one context could represent a pure technical

mean-reversion strategy on BTC/USD, while another could represent an AI-enhanced momentum strategy on a set of altcoins. Each context holds references to the modules it needs (perhaps a Technical Analyzer for a pure tech strategy, or both Technical + AITrendPredictor for an AI-backed strategy).

- **Unified Interface:** All strategies implement a common interface, say `StrategyBase`, with methods like `generate_signals()` and `update_positions()`. The `StrategyContext` calls these methods on schedule or on new data events, and gets back trade signals or recommendations. This uniform design allows the Execution Orchestrator to treat all strategies similarly when routing orders.
- **Independent Configuration:** Different strategies may have different config profiles (e.g., risk tolerance, what indicators or prompts to use, which symbols to trade). The config system supports **profile-based settings** so that each `StrategyContext` loads the appropriate parameters. For instance, one strategy might use a certain LLM prompt template and a different max cost setting, which it can obtain from the config.
- **Parallel Execution:** By leveraging Python's `asyncio` and task scheduling, multiple `StrategyContexts` can run in parallel without blocking each other. This is critical for a multi-symbol bot – each symbol/strategy combination is essentially an independent workflow, coordinated by the orchestrator. The system will be designed such that adding a new strategy is as simple as writing a new class (or plugin) and adding it to the config; the infrastructure (data feeds, execution, risk checks) is reused.

Overall, `StrategyContext` provides **modularity on the strategy level** – enabling the bot to run various strategies (traditional and AI) side by side, and making it easy to A/B test strategies or switch them on/off. It will also facilitate a potential “**strategy marketplace**” in future where strategies can be added or removed dynamically (a concept noted for future expansion ³⁹).

Execution Orchestrator (Decision Router & Broker Interface)

The **Execution Orchestrator** is the component that takes finalized trade decisions and turns them into reality on the exchanges. It serves two main purposes: deciding *if and how to execute* a signal, and interfacing with one or multiple broker/exchange APIs to place orders.

Key responsibilities:

- **Decision Consolidation:** The orchestrator receives trade decisions from strategies or the multi-agent Decision Engine. It will confirm that a decision has passed all risk checks and adjustments. In cases where multiple strategies produce signals for the same asset, it may need to reconcile them (for example, netting long vs short signals, or prioritizing one strategy over another based on a weight or a supervisor agent's input).
- **Order Routing & Multi-Broker Support:** The bot supports **multi-exchange trading** (e.g., Binance, Coinbase, OKX, Bybit) ⁴⁰. To manage this, the orchestrator uses an **Exchange Adapter** interface (e.g., via CCXT or custom adapters): each exchange has a module that translates generic order requests into API calls. The orchestrator can thus route an order to the appropriate exchange based on the symbol or a strategy's configuration. For instance, if a symbol is listed on multiple exchanges, the system could choose the one with the best liquidity or lowest fees. Multi-broker support also provides redundancy – if one exchange is down or unresponsive, the orchestrator can failover and execute on another (if the asset is available).
- **Order Management:** The orchestrator handles the full order lifecycle: placing orders (market, limit, etc. as determined by strategy), setting stop-loss/take-profit, monitoring order status, and confirming execution. It applies any final transforms needed (for example, splitting a large order into smaller chunks to reduce slippage).
- **Asynchronous Execution:** Using async I/O, the orchestrator can send orders to different exchanges in parallel and wait for confirmations without blocking the whole bot.

This is important when trading multiple symbols or when an exchange API is slow. - **Execution Policies:** It also implements execution logic policies. For example, the system might use a **two-tier execution** approach as noted earlier: if a very strong immediate signal is present (say from a fast indicator rule), the orchestrator might execute a small position immediately while the full AI analysis is still underway, then adjust or add to the position once the AI confirmation arrives ³ ⁴¹. This hybrid approach balances speed and intelligence. The orchestrator's policy module decides such cases based on signal confidence and urgency. - **Error Handling and Reporting:** If an order fails (API error, rejection, etc.), the orchestrator catches the exception and can retry or route to a backup exchange. All such events are logged. If a failure is critical (e.g., exchange not reachable), it can signal the RiskManager's emergency stop or notify the user.

The Execution Orchestrator is effectively the **gateway between the strategy world and the real market**. It will be thoroughly tested with sandbox/paper trading modes to ensure reliability. By abstracting away exchange specifics, we ensure the core strategies and AI logic remain exchange-agnostic. Adding a new broker in the future would mean writing a new adapter that implements a standard interface (`get_price()`, `place_order()`, etc.), without needing to alter the decision logic.

Configuration System

A robust **configuration system** is implemented to manage the many parameters and settings of this bot. It uses a combination of **YAML configuration files** for readability and easy editing, and **Pydantic** data models for validation and runtime access.

Structure of Configuration: - The primary config is a YAML file (or set of files) that define sections for each module: LLM settings, strategies, risk limits, exchange API keys, etc. This human-editable file allows quick tuning without touching code. - Pydantic `BaseModel` classes (e.g. `LLMConfig`, `StrategyConfig`, `RiskConfig`) mirror the structure of the YAML. When the bot starts, it parses the YAML into these models, providing type-checked Python objects throughout the system. This catches invalid or missing config values early (with clear error messages).

Key Configuration Features: - **Multi-Profile Support:** The system supports multiple config profiles (for example: `dev`, `paper`, `prod`). This can be handled either via multiple YAML files or a single file with nested profiles. The CLI (Typer) provides a `--profile` option to choose the configuration. For instance, a developer might use `--profile paper` to run the bot in paper trading mode with test API keys and safe limits, versus `--profile prod` for live trading. Profiles can inherit from each other to avoid duplication (common settings in a base, with overrides per profile). - **Secrets Management:** API keys and secrets (for exchanges and LLM APIs) are stored in the config but can be referenced via environment variables for security. Pydantic's use of `SecretStr` ensures they are handled carefully in code ⁴² ⁴³. The config loader will support `.env` files or environment variable overrides so that sensitive info isn't hardcoded in plain YAML. - **LLM Provider Chain & Model Selection:** The config defines which LLM provider to use as primary and can list fallback providers in order. It also contains specific model names per provider (e.g., GPT-4 vs GPT-3.5) and any model-specific parameters ⁸ ⁴⁴. The LLM Manager reads this to initialize accordingly. For local models, a path to the model file is configurable ⁴⁵. This approach makes it easy to switch providers or update model versions via config change. - **Cost and Performance Limits:** Configuration includes settings like `max_monthly_cost_usd` for LLM usage ⁴⁴, a toggle to enable/disable the AI components, and thresholds like acceptable latency. For instance, one could configure that if an LLM call takes longer than 10 seconds, the bot should proceed without it to avoid stalling – such knobs

can be exposed via config. - **Strategy & Risk Parameters:** Each strategy defined in the config can have its own sub-config (e.g., indicators to use, trading time window, max positions, etc.). Global risk limits such as max drawdown, max leverage, per-trade risk percent, VaR/CVaR limits, etc. are set here too. These are consumed by RiskManager and other modules to enforce policy. - **Hierarchical Fallback:** The system supports a hierarchy where if a specific setting for a strategy or profile is not set, it falls back to a default. For example, a general `RiskConfig` might set a global max drawdown, but a specific aggressive strategy profile could override it. If not, the global applies. This ensures **sensible defaults** are always in place even if not explicitly overridden (promoting safety).

The configuration is loaded at startup and can be hot-reloaded or validated via a CLI command. By using Pydantic, we ensure invalid config values (like a string where a number is expected, or a missing required field) will be caught immediately with clear error messages, preventing ambiguous runtime behavior. In summary, the config system provides **flexibility with guardrails** – any change to the bot's behavior goes through the config, which is version-controlled and type-checked.

Backtesting Engine (with PPO Training Support)

A critical part of the development and continuous improvement process is the **Backtesting Engine**. This module allows running the trading strategies (including the AI components) on historical data to evaluate performance, and it also provides a training environment for the PPO reinforcement learning module.

Design: - The backtester uses historical market data (candles, volume, maybe order books if available) for multiple symbols and simulates the bot's behavior over a specified period. It replays data in sequence (either bar by bar or lower timeframe for more granularity) and triggers the strategies and decision logic just like in live mode. - **Modular Data Input:** It can load data from CSV/Parquet files or via an API (for recent history). Using PyArrow to store the data ensures fast read/write and efficient slicing of large datasets. The data is provided to the Technical Analyzer and Sentiment Engine modules in the same format as live data, to ensure consistency.

Features: - **Multi-Symbol & Multi-Strategy Simulation:** The backtester can simulate all strategy contexts in parallel, similar to live, or focus on one at a time. It collects all trade decisions and executes them on a simulated exchange model. This model accounts for **slippage and latency** – e.g., if the strategy decided to buy at a price, the backtester can simulate a worse fill if volume was low (as per slippage estimation logic used in RiskManager) ³². It can also impose a delay between signal and execution to mimic real conditions (especially relevant if waiting for LLM). - **PPO Training Mode:** In a special mode, the backtester integrates with the PPO Adjuster to train its policy. Essentially, we treat each backtest run (or multiple runs) as episodes in an RL environment. The state is the market and portfolio state, actions are adjustments (or can even be the decision to trade or not), and reward is computed from trade outcomes with risk considerations (e.g., a trade's profit minus a penalty for large drawdowns). We leverage PPO algorithms (via stable-baselines3 or custom code) to update the policy network. This training can be iterative: run backtest -> update model -> repeat, until some convergence or performance is achieved. The result is a trained model that can be saved and later loaded by the live bot for the PPO Adjuster. (This process is somewhat experimental, so it will be initially done offline and carefully evaluated.) - **Metrics and Reporting:** The backtester will output detailed metrics: win rate, P/L, drawdown, Sharpe ratio, max VaR, etc. for each strategy both with and without AI enhancements. This allows us to quantify the benefit of the AI layer and PPO. The design target is that sentiment and LLM should add a few percentage points to win rate in backtests ⁴⁶, and PPO should improve risk-adjusted returns (e.g. Sharpe up by 0.1-0.2) ⁴⁷. If not, these

components need tweaking before live use. - **CLI Integration:** We will provide CLI commands (via Typer) like

```
bot backtest --symbol BTC --strategy AITrend --start 2021-01-01 --end 2022-01-01
```

to easily run simulations. Results can be saved to files or printed. We may also integrate with Jupyter notebooks for analysis and visualization of trades.

The backtesting engine not only serves for initial development but will be used continuously to test strategy updates. It enables **safe experimentation**: any new change (new model version, new indicator, parameter tweak) should be backtested and perhaps even paper-traded before going live with real money.

Monitoring and Observability

Observability is crucial for a trading bot that may run 24/7. We incorporate multiple layers of monitoring, logging, and alerting to ensure the system's behavior is transparent and any issues are caught early.

- **Structured Logging:** All components will use Python's logging facilities to output structured logs (e.g., in JSON) including timestamps, module names, and context (symbol, strategy IDs, etc.). Important events like trade decisions, orders placed, errors, or unusual AI outputs are logged at appropriate levels. This allows later analysis or debugging. Logs can be routed to both console and files. In production (GCP or similar), logs would be collected by a centralized logging service.
- **Metrics & Prometheus:** The bot will expose key metrics via **Prometheus**. A background task or the FastAPI app will periodically update metrics such as: current portfolio value, PnL, number of trades, win rate, latency of LLM calls, API call counts, etc. These metrics can be scraped by Prometheus and visualized in Grafana dashboards ⁴⁸. For example, we can have a dashboard panel showing the frequency of emergency stops or the average confidence of AI predictions over time.
- **Real-Time Dashboard:** Optionally, a web dashboard (React frontend or even just Grafana UI) will display real-time status: open positions, performance graphs, recent signals, and system health (API latency, LLM usage vs. monthly budget, etc.). This helps the user or developer to understand what the bot is doing at a glance.
- **Alerts (Telegram/Email):** Critical events trigger instant alerts. The system integrates with a Telegram bot API (and/or email/SMS) to send notifications. For instance, if an **emergency stop** is triggered and trading halts, an alert with the reason is sent immediately ³⁸. Other examples: if the LLM monthly cost is nearing the limit, send an alert; if a strategy has consecutive losses beyond a threshold, notify to review; or simply periodic daily summaries of performance pushed to the user's phone. These thresholds and channels are configurable.
- **Health Checks:** The FastAPI service can provide health endpoints (e.g. `GET /health`) that return status of sub-systems (connected to exchange? LLM responding? etc.). In GCP deployment, this can be used for automated monitoring (Cloud Watch or GKE health checks) to restart components if hung. Internally, the bot can also self-monitor: a watchdog coroutine can ensure the main event loop is running and no part has stalled. If it detects an anomaly (e.g., no trade or no heartbeat for a while), it could restart certain sub-systems or at least log and alert.
- **Audit Trail & Explainability:** Every trade decision made by the AI components comes with a "reasoning" (thanks to the LLM's JSON output which includes bullet point rationale). We log this reasoning with the trade. This is invaluable for later analysis or for user trust – one can trace *why* the bot made a certain trade (e.g., "LLM indicated bullish due to rising RSI and positive news sentiment, confidence 70%"). Even though an AI is in the loop, we maintain as much transparency as possible into its decision process.

In summary, the bot will not run as a black box. Comprehensive monitoring ensures we know its state and can be confident in its actions. Safe defaults extend here as well: for example, if monitoring detects unusual activity (like a sudden flood of trade signals), the system could automatically pause and request human review. The emphasis is on **operational safety and clarity** from day one of deployment.

Implementation Plan

Implementing this system will be done in **iterative phases** to manage complexity and allow incremental testing. We outline a plan with phases, each roughly 2 weeks (some shorter, some longer), totaling around 18-20 weeks of development. Each phase produces a working subset of the system that can be tested before moving on. Below is the plan with key tasks, dependencies, and estimated durations:

Phase 1: Foundation and Core Trading Engine (Weeks 1-3)

Objectives: Establish the basic multi-symbol trading framework, real-time data ingestion, and simple execution on a single exchange. Implement critical risk safety features from the start.

Tasks: - Setup project scaffolding, choose libraries, and write basic **Exchange Adapter** for a primary exchange (e.g., Binance spot). Include robust error handling for API calls ⁴⁹. - Implement market data streaming or polling for price updates (using exchange REST APIs or websockets). Begin storing data in memory (possibly using PyArrow for efficiency). - Develop the **Logging infrastructure** (structured logs with unique IDs for trades). - Implement a basic **RiskManager** with checks: order book depth slippage estimation, liquidity limits, and a simple global stop-loss (e.g., if portfolio loss > X%) ³² ³³. - Implement an **EmergencyStopSystem** that can cancel all orders and halt trading on a trigger (with a manual override and automated triggers as needed) ⁵⁰ ³⁷. - Create a simple **Technical Strategy** (e.g., moving average crossover or RSI-based) as a placeholder strategy for testing the pipeline. - Build the **Execution Orchestrator** with the ability to place market orders through the exchange adapter and handle order responses.

Dependencies: Basic knowledge of exchange API and available data; no external AI components yet, so this can proceed independently.

Deliverables & Success Criteria: By end of Phase 1, the bot can connect to the exchange, retrieve price data for multiple symbols, generate a basic trade signal, and execute a trade (in a test or paper environment) obeying risk checks. It should run continuously without crashing for a period (e.g., 24-48h in paper trade). All critical failures (lost connection, etc.) should be caught and logged or handled gracefully.

Phase 2: Traditional Strategy Expansion & Backtesting Framework (Weeks 4-5)

Objectives: Build out the Technical Analyzer and backtesting capabilities. Establish baseline performance of a purely technical bot as a benchmark.

Tasks: - Implement the **Technical Analyzer** fully: integrate TA-Lib (or equivalent) and calculate key indicators (RSI, MACD, Bollingers, MAs, ATR, etc.) for incoming data. Ensure it supports multi-timeframe data. - Develop multiple simple **StrategyContext** instances using the Technical Analyzer outputs (e.g., a momentum strategy and a mean-reversion strategy) to test multi-strategy support. - Introduce the

Backtesting Engine: be able to run a simulation on historical data for a given strategy. Include modeling of slippage and trade execution delay in backtests to mimic real conditions ⁵¹. - **Paper Trading Mode:** Extend the execution layer to support a simulated broker (paper trading), which the backtester will also use. This could be as simple as tracking an internal portfolio state when orders are "executed". - Validate and refine the RiskManager rules with backtest data. For example, verify that liquidity and slippage checks catch unrealistic trades. - Start a **documentation log** of strategy performance: record the win rate, max drawdown, Sharpe, etc., for these baseline strategies via backtesting.

Dependencies: Completion of Phase 1 (core infrastructure and basic strategy execution) is required. TA-Lib (or `ta`) needs to be installed and data accessible.

Deliverables: A working backtester that can replay historical data and output performance metrics. At least one traditional strategy configured and yielding a realistic ~50% win rate in backtests with drawdown under control (e.g., <25%) ⁵². This provides a baseline to improve upon with AI. By this point, the project should have solid core functionality and be ready to incorporate AI features.

Phase 3: LLM Integration Layer (Weeks 6–7)

Objectives: Introduce AI capabilities by integrating the LLM Manager and hooking it into the trading logic for predictions.

Tasks: - Implement the **LLMManager** class with multi-provider support, cost tracking, and caching as per design ⁷. Begin with OpenAI GPT-4 as primary (if available) and perhaps a local model for testing. Include the fallback chain logic and basic response validation. - Define the initial **system prompt and output format** for trading predictions (as JSON). Leverage the prompt guidelines (e.g., instruct the LLM to use only provided data, output JSON, etc.) ²³ ²⁴. - Build a minimal **AITrendPredictor** that queries the LLM for a single timeframe prediction (to start). Feed it a summary of technical indicators for the symbol. This will be a simplified version to prove the concept. - Implement caching of LLM responses for e.g. 1 hour for each symbol's prediction to control costs ⁷. - **Cost Monitoring:** Add to the monitoring module a way to accumulate API costs (perhaps parse OpenAI billing if possible, or estimate based on tokens) and enforce monthly limits (the bot can stop AI calls or switch to a cheaper model if the budget is nearly hit). - Write unit tests for LLMManager using a dummy model (to simulate various failures and ensure fallback works).

Dependencies: Phase 2's data and strategy framework is needed so that we have technical data to feed into the LLM. Also need API keys and access for LLM providers (setup in config).

Deliverables: The bot can now produce an AI-based prediction for a symbol on request. We should be able to log something like: *"AI predicts BTC/USD bullish 60% confidence, expected +3% move, reasoning [...]"*. The multi-provider setup should be tested by simulating a failure of the primary and seeing the fallback engage ⁵ ⁹. By end of Phase 3, the groundwork for AI assistance is laid, though it may not yet be fully integrated into live trading decisions.

Phase 4: Multi-Modal Sentiment & News Integration (Weeks 8–9)

Objectives: Enrich the AI's input by adding the Sentiment Engine and news data, and incorporate these into the AITrendPredictor for a multi-modal analysis.

Tasks: - Implement the **SentimentEngine** to fetch data from Reddit (via PRAW or another API) and parse a few crypto subreddits for mention counts and sentiment. Also implement a lightweight Twitter sentiment fetch (if API access is available; if not, consider using an alternative or skipping Twitter due to cost) ¹² . - Integrate free **news feeds**: use an RSS feed parser to gather recent article headlines for the top news sites ¹³ . Implement a simple sentiment scoring for headlines (e.g., count positive vs negative words) or use an LLM in a prompt to summarize news sentiment if feasible. - Pull **Fear & Greed Index** from a public source (it's usually updated daily) ¹⁴ and any other free sentiment indicators (like alternative.me API). - Combine these in `SentimentData` - design a schema with fields like overall sentiment score, twitter_score, reddit_score, news_score, fear_greed, etc. Implement an `aggregate()` method to compute a weighted overall sentiment ¹⁶ . - Update the **AITrendPredictor** prompt building to include sentiment info and top news headlines ²¹ ⁵³ . Also include any important market context (e.g., overall market cap or dominance, if relevant). - Test the prediction end-to-end: does the inclusion of sentiment change the outputs meaningfully? Backtest the AI-augmented strategy on recent data and see if signals align with actual market moves (qualitatively). - Monitor the latency and cost impact: ensure that adding sentiment (with multiple API calls) doesn't slow the system unacceptably. Use async calls to fetch Reddit/news in parallel.

Dependencies: Needs Phase 3 (LLM integration) complete, since the LLM will be interpreting the sentiment/news data. Also requires internet access for sentiment sources or cached data for testing.

Deliverables: A more comprehensive AITrendPredictor that uses multi-modal data. We should see in logs that the prompt now contains sentiment and news context, and the LLM's reasoning perhaps references these ("Bullish sentiment on Reddit and positive news about ETF approval..."). Backtests or paper tests should show a slight performance uptick with sentiment (target +2-5% win rate in simulations) ⁴⁶ , validating the feature. The system should still meet the performance budget: ideally AI inference per symbol per timeframe stays within ~5 seconds on average (with caching), so that we can handle multiple symbols without lag.

Phase 5: Multi-Agent Decision Making System (Weeks 10-12)

Objectives: Scale the single LLM prediction into a **multi-agent framework** to improve decision robustness, and integrate it fully into the live trading decision flow.

Tasks: - Define and implement the various **Agent classes**: e.g., `AnalystAgent` (the primary AITrendPredictor's role), `RiskAgent` (could be a rule-based or LLM-based agent focusing purely on risk evaluation), `DevilsAdvocateAgent` (an LLM agent that challenges the Analyst's decision) ⁵⁴ ⁵⁵ , and a `SupervisorAgent` (can be a simple rule-based overseer that has veto power if agents disagree strongly or if risk is too high). - Develop the **Decision Fusion Engine** that orchestrates these agents. For example: AnalystAgent provides a draft trade decision (buy/sell/hold with confidence and reasoning), the RiskAgent reviews risk aspects (maybe using a different prompt or just the RiskManager's outputs), the Devil's AdvocateAgent tries to find contrary evidence ⁵⁶ ⁵⁷ , and then the SupervisorAgent decides to approve, modify, or reject the trade. This could be implemented by sequential async calls or concurrently if feasible. - Ensure that the multi-agent process still returns a decision in a reasonable time (the target was <30 seconds for full deliberation) ⁵⁸ . This might involve using faster, smaller models for some agents or simplifying their tasks. - Integrate the final decision from the multi-agent engine into the Execution Orchestrator. At this point, the orchestrator should wait for the AI/multi-agent signal (unless a fast-path rule already triggered a small trade). Implement logic to handle cases where agents disagree - possibly require consensus or at least Supervisor approval for any action to proceed. - **Precompute & Parallelize:** If monitoring multiple

symbols, consider precomputing some LLM predictions in the background. For example, maintain a rolling analysis of a watchlist so that when a signal condition triggers, the AI insight is already cached (to avoid long waits). Use `asyncio.create_task` to continuously update AI predictions for each symbol every N minutes in the background. - Test the multi-agent system in paper trading: intentionally create scenarios for the Devil's Advocate (e.g., the Analyst says buy but fabricate a risk factor to see if DA catches it) to ensure the logic works. Also test that the Supervisor correctly halts trades when needed (e.g., if RiskAgent flags a major risk, no trade should happen).

Dependencies: Requires Phase 4 fully working (AI predictions with sentiment) because agents will rely on those capabilities. Also needs RiskManager outputs to feed RiskAgent.

Deliverables: The trading decisions are now the result of a coordinated multi-agent process, adding an extra layer of scrutiny. We expect more **consistent and safer decisions** – e.g., fewer false positives because the Devil's Advocate might filter out over-optimistic trades, and overall system behaving more like a team of experts confirming a trade ⁵⁹. By end of Phase 5, the bot in paper trading should achieve roughly ~50% win rate with reduced catastrophic trades (none that violate risk rules), meeting a key benchmark ⁵⁸. It should run for a couple of weeks in a paper account without major issues, demonstrating stability.

Phase 6: Risk-Aware PPO and Continuous Learning (Weeks 13–15)

Objectives: Implement the PPO reinforcement learning adjuster and train/tune it using the backtesting framework. Integrate its adjustments into live decisions.

Tasks: - Develop the PPO model architecture (could start with a simple neural network with inputs for key state metrics: recent returns, volatility, VaR, etc., and output a factor to adjust position size or a bias to adjust LLM confidence). - Set up the **training loop** within the backtester: define the reward function (e.g., trade outcome in terms of risk-adjusted profit, penalize exceeding risk thresholds or big drawdowns). Use historical data to simulate many episodes. This may require generating many random strategy scenarios or using a long history and having the agent slowly learn. - Use a stable RL library or implement PPO update steps (policy and value network optimization) if doing from scratch. Ensure to tune hyperparameters and possibly do early stopping if the policy converges ⁶⁰. - Once a reasonable policy is learned (not necessarily super optimal, but at least not harmful), integrate the **RiskAwarePPO** adjuster into the live decision flow. This means when a trade decision comes out of the multi-agent engine, pass it (and current state) to the PPO agent's `adjust_decision()` method which returns adjusted position size or modified action recommendation ³⁰. - Implement safeguards: if the PPO suggests something extreme or contradictory (being an ML model, one must be careful), have the RiskManager constrain it. For example, if PPO output would increase position beyond allowed max, cap it. - Begin with PPO in advisory mode (maybe log what it *would* do, but still follow the original plan) for a testing period. This will allow comparing performance with/without PPO adjustments. - Gradually start using PPO adjustments in paper trading and measure impact. The expectation is a modest improvement in Sharpe or reduction in volatility of returns ⁴⁷. - **Continuous Learning Pipeline (optional):** If feasible, set up a process to periodically retrain or update the PPO model with the latest data (perhaps monthly). This could be automated but with human oversight (to avoid the bot drifting unpredictably).

Dependencies: Phase 2's backtester must be functional to train PPO. Also, we assume by now we have accumulated enough trade simulation data to train on. Multi-agent system (Phase 5) should be done, as PPO will adjust its output.

Deliverables: A trained PPO model integrated into the runtime. By end of Phase 6, tests should show that when PPO adjuster is on, *undesirable trades are further downsized or skipped*, improving the risk profile. For example, perhaps the maximum drawdown in a backtest improves a few percentage points with PPO active. This phase is marked optional because if the PPO does not prove beneficial in testing, we may choose to omit or delay it for production. Nonetheless, the infrastructure for it will be in place for further experimentation.

Phase 7: User Interface, Monitoring & Final Refinements (Weeks 16–17)

Objectives: Build out the user-facing aspects (web dashboard or CLI enhancements), finalize monitoring, and prepare the system for a production-like environment.

Tasks: - Develop a minimal **web dashboard** using FastAPI and a simple React frontend (if time permits). At minimum, create API endpoints for retrieving current status: open positions, recent trades, performance stats, etc., and display these in a web page or Grafana. If not React, at least set up Grafana with Prometheus as planned and create a couple of dashboards for internal use. - Implement **Telegram bot integration** for alerts and possibly basic commands. For example, on certain triggers (emergency stop, large loss, etc.), send a Telegram message ³⁸. Also allow the bot to receive a manual “pause” or “resume” command via Telegram for convenience. - Harden the logging and config: rotate logs, archive old logs, ensure no sensitive info is logged. Test config profile switching to make sure switching from paper to live is as simple as changing the profile name and restarting. - Security audit: ensure API keys are encrypted or at least not exposed (use environment variables, etc.), lock down the FastAPI endpoint (if remote access is enabled, use authentication). - **Performance tuning:** Profile the system for any bottlenecks (CPU usage, memory). For example, if using a local LLM model, ensure it runs on GPU if available; if sentiment fetch is slow, consider threading or caching more. Aim to keep the system efficient enough to run on a single machine or VM. - Write **documentation** for usage: how to configure, how to add a new strategy, how to interpret logs and outputs. Also document how each module works for future maintainers.

Dependencies: The core system must be feature-complete by now (phases 1-6). This phase focuses on polish and usability.

Deliverables: By the end of Phase 7, the bot is essentially ready to go live. We should have a user-friendly way to monitor it (web or at least comprehensive logs/metrics) and a clear understanding of its performance. All modules are integrated, and the team (or user) is comfortable that the bot can be left running with alerts set up for any important event.

Phase 8: Testing, Deployment and Gradual Rollout (Weeks 18–20)

Objectives: Rigorously test the system end-to-end, deploy it to production environment (e.g., GCP), and start live trading with small capital in a controlled manner.

Tasks: - **End-to-End Testing:** Conduct various tests: unit tests (should be mostly done per phase), integration tests (simulate a day of trading with recorded real data), and stress tests (run with many symbols or with artificial high-frequency data to see where it breaks). Also test failure scenarios: e.g., simulate an exchange API outage and ensure the bot’s emergency stop triggers properly. - **Deployment Setup:** Containerize the application using Docker. Write a Dockerfile that sets up Python environment, etc. Then use **Infrastructure-as-Code** (possibly Terraform or just manual setup scripts) to deploy on GCP. The

target could be a VM or a Kubernetes cluster for high availability. Set up necessary environment on the server (docker-compose with Prometheus/Grafana if needed, volumes for persistent data, etc.) ⁶¹ . - **Monitoring in Prod:** Deploy Prometheus and Grafana on GCP (or use a managed service) and point it to our bot's metrics endpoint. Verify that alerts (Telegram/email) work from the cloud environment. - **Gradual Capital Deployment:** Start live trading in a sandbox or with a very small amount (e.g. \$100 as suggested) ⁶² . Let it run for a week. Observe behavior, and manually intervene if something looks off. If performance is acceptable and no major bugs, increase to \$500, then further gradually. This phased ramp-up ensures if any bug or unexpected behavior occurs, losses are minimized. - Review the live trading logs and metrics daily. Check that the live performance aligns with backtest expectations. Pay attention to any discrepancies (could indicate issues like data feed differences or model overfitting). - Finalize any fixes or tweaks discovered during this rollout. For example, if latency of AI is a problem, maybe reduce frequency of AI calls or further utilize caching.

Dependencies: All prior phases completed. Need access to production accounts, cloud resources, etc.

Deliverables: The bot running in a production environment, trading with real money cautiously. By end of week 20, the goal is to have a stable system with perhaps a small positive return or at least a break-even (after fees) performance, demonstrating that the AI enhancements did not introduce losses. At this point, the architecture supports adding more exchanges or scaling up capital. Ongoing work would involve strategy refinement, continuous learning, and maintenance, but the core implementation is complete.

Throughout these phases, **safe defaults and iterative testing** guide development. Early phases establish safety (risk checks, kill-switches) so that later AI experiments never threaten the stability of the bot. By following this plan, we incrementally build a sophisticated AI-enhanced trading system with a strong foundation in traditional trading principles and modern AI techniques. Each component is modular and testable, aligning with the goal of high extensibility and maintainability for the long term.

Sources: The design and plan were informed by the expert review document and AI enhancement proposal, ensuring we integrated recommended best practices (multi-agent consensus, VaR/CVaR risk management, cost-effective sentiment data, etc.) and addressed prior concerns (latency, slippage, data quality) ⁶³ ¹² . This comprehensive approach aims to deliver an advanced trading bot that is both powerful and safe.

¹ ² ⁵ ⁶ ⁸ ⁹ ¹⁷ ²⁰ ²¹ ²² ²⁵ ²⁶ ²⁷ ²⁸ ²⁹ ³⁰ ³⁹ ⁴⁰ ⁴² ⁴³ ⁴⁴ ⁴⁵ ⁵³ per_plexity.pdf

file:///file_000000005bb061f7b806181939b55c53

³ ⁴ ⁷ ¹⁰ ¹¹ ¹² ¹³ ¹⁴ ¹⁵ ¹⁶ ¹⁸ ¹⁹ ²³ ²⁴ ³¹ ³² ³³ ³⁴ ³⁵ ³⁶ ³⁷ ³⁸ ⁴¹ ⁴⁶ ⁴⁷ ⁴⁸ ⁴⁹ ⁵⁰ ⁵¹ ⁵²
⁵⁴ ⁵⁵ ⁵⁶ ⁵⁷ ⁵⁸ ⁵⁹ ⁶⁰ ⁶¹ ⁶² ⁶³ ai-trading-bot-review.md

file:///file_0000000027b4622f838b2a0014124177